

Host-Based Intrusion Detection System Using the ELK Stack and Suricata

Cybersecurity Project

PREPARED BY

ATHUL KRISHNA A U

CONTENTS

INTRODUCTION	1
METHODOLOGY	3
INSTALLATION	4
CONFIGURATION.....	6
ELASTICSEARCH.....	6
KIBANA.....	8
LOGSTASH.....	13
SURICATA	16
FILEBEAT	18
DASHBOARD SETUP.....	22
CONCLUSION	24

Introduction

In this project, I am going to make a **HIDS** by configuring Suricata, VirusTotal , AlienVault OTX and ELK Stack together to analyse real-time logs. So let's familiarize with it.

1) What is HIDS?

It is a cybersecurity tool that monitors and analyses traffic to host for suspicious activities, policy violations, or potential attacks in real-time.

If malicious activity is detected, the HIDS generates alerts for security analysts to investigate.

2) What is ELK Stack?

The ELK Stack is a powerful open-source log management and data analysis platform used for collecting, storing, analysing and visualizing logs or data from multiple sources.

ELK stands for:

- **Elasticsearch:** A search and analytic engine that stores and indexes data for fast searching.
- **Logstash:** A data processing pipeline that collects, filters, and sends data to Elasticsearch.
- **Kibana:** A visualization tool that displays data stored in Elasticsearch using dashboards and graphs.

3) What is VirusTotal API?

The VirusTotal API is a RESTful API provided by VirusTotal that allows tools and systems like Elastic Stack or Suricata-based HIDS setup to automatically query VirusTotal's database for threat intelligence information

4) What is AlienVault OTX?

AlienVault OTX (Open Threat Exchange) is a threat intelligence sharing platform by AT&T Cybersecurity. It provides Indicators of Compromise (IOCs) — such as:

- Malicious Ips
- Domains
- URLs
- File hashes
- Threat actor information.

You can access this data through the OTX API using your API key.

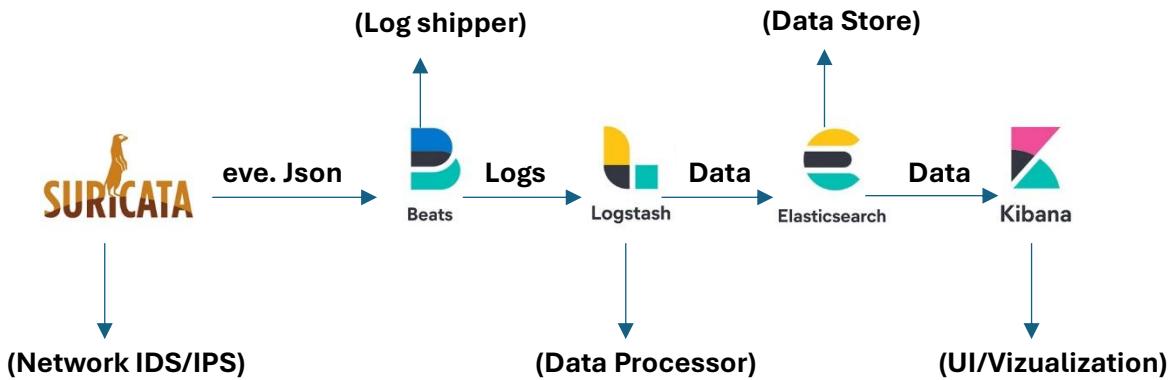
5) What is Suricata?

Suricata is an open-source Intrusion Detection system (IDS), Intrusion Prevention System (IPS) and Network Security Monitoring (NSM) tool developed by Open Information Security Foundation (OISF).

It monitors network traffic in real-time and analyses packets to detect malicious activities, attacks, or policy violations.

Suricata works by capturing packets from a network interface and inspect them using **rule-based signatures**. It then identifies attacks or anomalies and logs them for analysis.

Suricata generates the **eve.json** file through its **EVE (Extensible Event Format) logging module**, which is designed to output structured, machine-readable logs for use in analysis platforms like **ELK Stack**, **Splunk** or other SIEM tools.



- Suricata writes EVE JSON logs to disk.
- Filebeat monitors these logs, parses them using the Suricata module, and sends them to Logstash (or directly to Elasticsearch).
- Logstash (if used) processes and enriches the logs.
- Elasticsearch stores and indexes the processed logs.
- Kibana queries Elasticsearch to visualize and analyse the Suricata event data.

Methodology

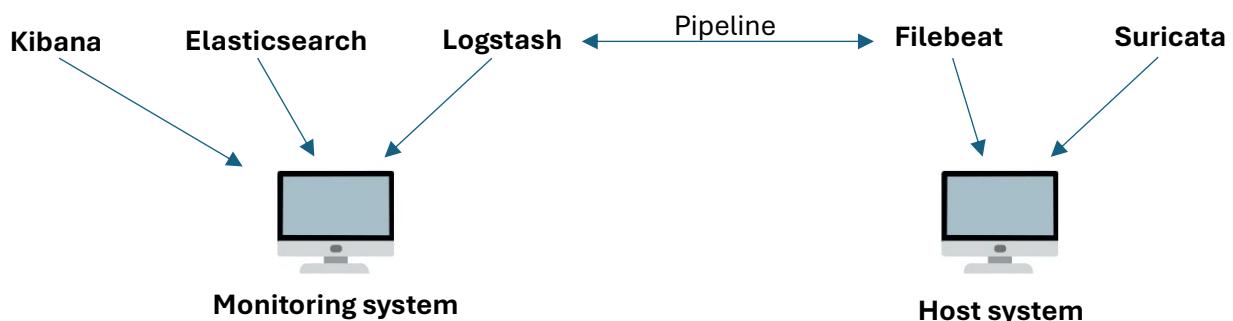
1) Tools required:

- Elasticsearch
- Logstash
- Kibana
- Filebeat
- Suricata
- VirusTotal
- AlienVault OTX

2) OS used:

Here I am using **Debian based OS (Kali Linux)**. Any other OS such as Windows, MacOS, etc can also be used. Many other Linux-based operating systems, including Ubuntu, Mint, and Kali Linux, are based on Debian.

Both host and monitoring system are set on kali linux.



Installation

1) Elasticsearch

Download Elasticsearch package from [elastic.co](https://www.elastic.co)

<https://www.elastic.co/downloads/elasticsearch>

Here I am using kali linux so I am going with **deb x86_64** package.

Install Elasticsearch:

`dpkg -i package_name.deb`

for eg. `dpkg -i elasticsearch-9.1.5-amd64.deb`

Start the Elasticsearch service:

`sudo systemctl enable elasticsearch`

`sudo systemctl start elasticsearch`

To check the elasticsearch service status:

`sudo systemctl status elasticsearch`

2) Logstash

Download Logstash.

<https://www.elastic.co/downloads/logstash>

Install logstash:

`dpkg -i package_name.deb`

Start Logstash service:

`sudo systemctl enable logstash`

`sudo systemctl start logstash`

`sudo systemctl status logstash`

3) Kibana

Download Kibana package:

<https://www.elastic.co/downloads/kibana>

Install Kibana:

`dpkg -i package_name.deb`

Start Kibana service:

`sudo systemctl enable kibana`

`sudo systemctl start kibana`

```
sudo systemctl status kibana
```

Next download and install Filebeat and Suricata in host system.

4) Filebeat

Download Filebeat package:

```
https://www.elastic.co/downloads/beats/filebeat
```

Install Filebeat:

```
dpkg -i package_name.deb
```

Start Filebeat service:

```
sudo systemctl enable filebeat  
sudo systemctl start filebeat  
sudo systemctl status filebeat
```

5) Suricata

Add the OISF Repository:

```
sudo add-apt-repository ppa:oisf/suricata-stable
```

Update Package Lists:

```
sudo apt update
```

Install Suricata:

```
sudo apt install suricata
```

Start and Enable Suricata

```
sudo systemctl enable suricata  
sudo systemctl start suricata  
sudo systemctl status suricata
```

Next step is to configure these tools together by editing their yml/yaml, conf, rules file.

Note:

Don't forget to restart the service after editing their file if the services are running.

For eg. if you edited the elasticsearch.yml file then restart the service by

```
sudo systemctl restart elasticsearch
```

Configuration

1) Configure Elasticsearch

Go to `/etc/elasticsearch` directory and edit the `elasticsearch.yml`

```
sudo nano elasticsearch.yml
```

Uncomment required lines

`Network.host: "0.0.0.0"` #using 0.0.0.0 exposes the service to all network connections.

`http.port: 9200`

`transport.host: "0.0.0.0"`

```
#----- Paths -----
# Path to directory where to store the data (separate multiple locations by comma):
#path.data: /var/lib/elasticsearch
# Path to log files:
#path.logs: /var/log/elasticsearch
#----- Memory -----
# Lock the memory on startup:
#bootstrap.memory_lock: true
# Make sure that the heap size is set to about half the memory available
# on the system and that the owner of the process is allowed to use this
# limit.
# Elasticsearch performs poorly when the system is swapping the memory.
#----- Network -----
# By default Elasticsearch is only accessible on localhost. Set a different
# address here to expose this node on the network:
network.host: "0.0.0.0"
# By default Elasticsearch listens for HTTP traffic on the first free port it
# finds starting at 9200. Set a specific HTTP port here:
http.port: 9200
# For more information, consult the network module documentation.
#----- Discovery -----
```

```
# Enable security features
xpack.security.enabled: true
xpack.security.enrollment.enabled: true

# Enable encryption for HTTP API client connections, such as Kibana, Logstash, and Agents
xpack.security.http.ssl:
  enabled: true
  keystore.path: certs/http.p12

# Enable encryption and mutual authentication between cluster nodes
xpack.security.transport.ssl:
  enabled: true
  verification_mode: certificate
  keystore.path: certs/transport.p12
  truststore.path: certs/transport.p12
# Create a new cluster with the current node only
# Additional nodes can still join the cluster later
cluster.initial_master_nodes: ["RA-7"]

# Allow HTTP API connections from anywhere
# Connections are encrypted and require user authentication
http.host: "0.0.0.0"

# Allow other nodes to join the cluster from anywhere
# Connections are encrypted and mutually authenticated
transport.host: "0.0.0.0"

#----- END SECURITY AUTO CONFIGURATION -----
```

Save and exit:

Ctrl+x

y, and enter

Enable and start the elasticsearch service, if already running then restart it.

```
sudo systemctl restart elasticsearch
```

If you want to restset the elastic password, then run this command

```
sudo /usr/share/elasticsearch/bin/elasticsearch-reset-password -u elastic
```

If you get a permission error then run these commands

```
sudo chown -R elasticsearch:elasticsearch /etc/elasticsearch
```

```
sudo chown -R 750 /etc/elasticsearch
```

Then retry:

```
sudo /usr/share/elasticsearch/bin/elasticsearch-reset-password -u elastic
```

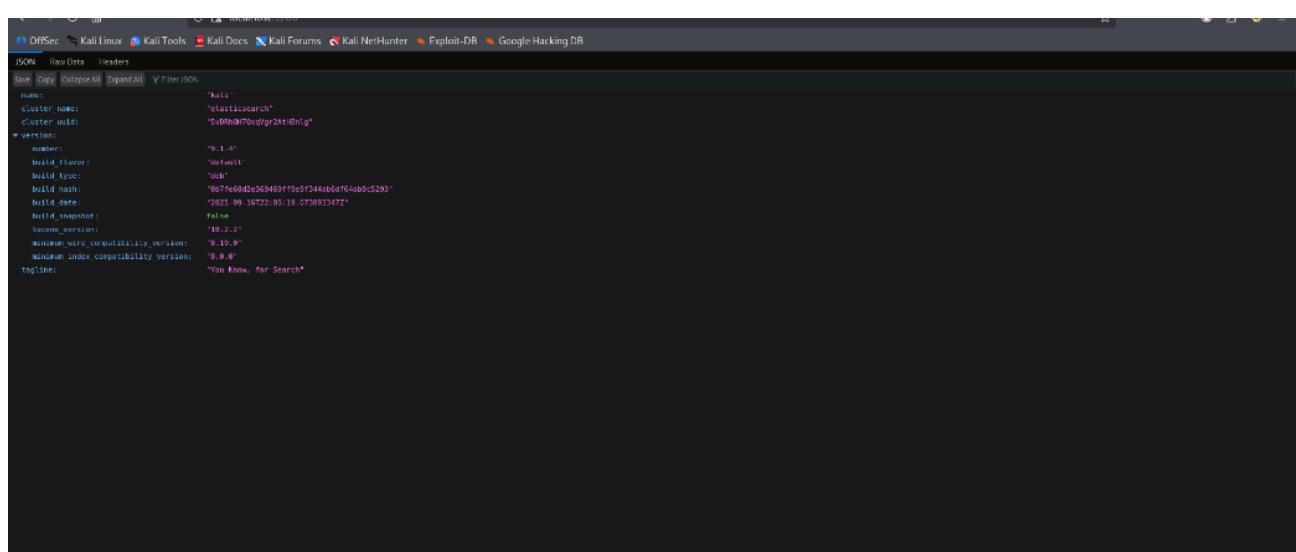
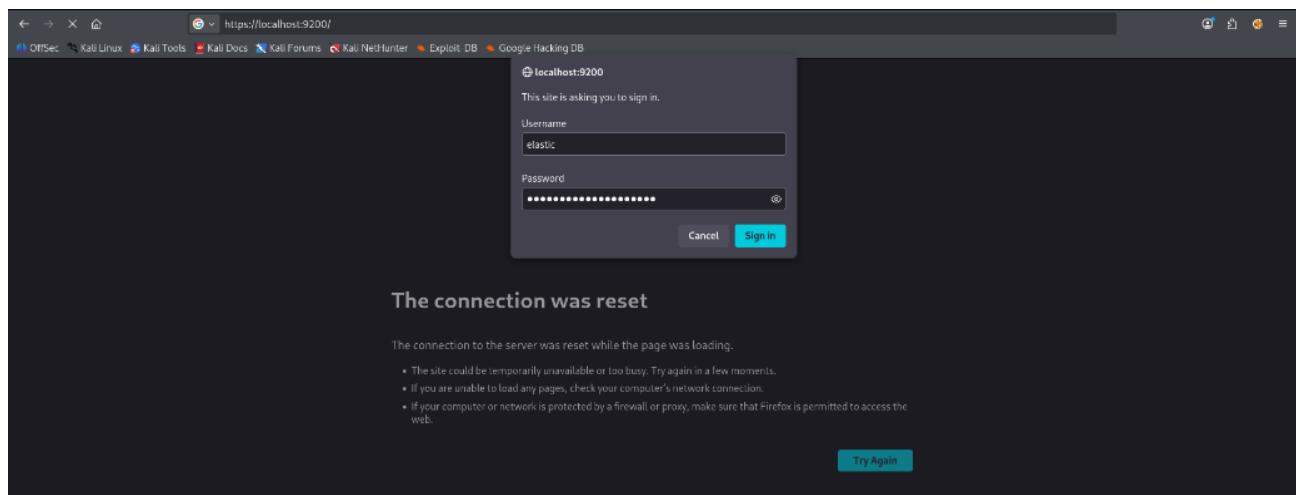
```
L:~# sudo /usr/share/elasticsearch/bin/elasticsearch-reset-password -u elastic --url https://localhost:9200
This will reset the password of the [elastic] user to an autogenerated value.
The password will be printed in the command output.
Please confirm that you would like to continue [y/N]

[password for the elastic user successfully reset.
New value: P7cIqG0AMLi4z9k1ch]
```

Note it somewhere, don't forget the password.

Check the elasticsearch is working or not:

Go to your browser and <https://localhost:9200>



Your Elasticsearch instance is running perfectly if you got a result like this.

2) Configure Kibana

Go to `/etc/kibana` directory and run this command

```
sudo nano kibana.yml
```

Uncomment these lines in **System: Kibana server** section

```
server.port: 5061
```

```
server.host: "0.0.0.0"
```

```
GNU nano 8.6                               /etc/kibana/kibana.yml *
# For more configuration options see the configuration guide for Kibana in
# https://www.elastic.co/guide/index.html

# ===== System: Kibana Server =====
# Kibana is served by a back end server. This setting specifies the port to use.
server.port: 5061

# Specifies the address to which the Kibana server will bind. IP addresses and host names are both valid values.
# The default is 'localhost', which usually means remote machines will not be able to connect.
# To allow connections from remote users, set this parameter to a non-loopback address.
server.host: "0.0.0.0"

# Enables you to specify a path to mount Kibana at if you are running behind a proxy.
# Use the 'server.rewriteBasePath' setting to tell Kibana if it should remove the basePath
# from requests it receives, and to prevent a deprecation warning at startup.
# This setting cannot end in a slash.
#serverbasePath: ""

# Specifies whether Kibana should rewrite requests that are prefixed with
# 'serverbasePath' or require that they are rewritten by your reverse proxy.
# Defaults to 'false'.
#server.rewriteBasePath: false

# Specifies the public URL at which Kibana is available for end users. If
# 'serverbasePath' is configured this URL should end with the same basePath.
#server.publicBaseUrl: ""

# The maximum payload size in bytes for incoming server requests.
#server.maxPayload: 1048576

# The Kibana server's name. This is used for display purposes.
server.name: "Kibana"

# ===== System: Kibana Server (Optional) =====
# Enables SSL and paths to the PEM-format SSL certificate and SSL key files, respectively.

^G Help          ^O Write Out      ^F Where Is      ^K Cut           ^T Execute        ^C Location       M-U Undo        M-A Set Mark     M-] To Bracket
^X Exit         ^R Read File      ^\ Replace       ^U Paste         ^J Justify        ^/ Go To Line    M-E Redo        M-6 Copy        ^B Where Was
```

In **System: Elasticsearch** Section, give `elasticsearch.hosts`, elasticsearch's username and password.

```
# ===== System: Elasticsearch =====
# The URLs of the Elasticsearch instances to use for all your queries.
elasticsearch.hosts: ["https://localhost:9200"]

# If your Elasticsearch is protected with basic authentication, these settings provide
# the username and password that the Kibana server uses to perform maintenance on the Kibana
# index at startup. Your Kibana users still need to authenticate with Elasticsearch, which
# is proxied through the Kibana server.
elasticsearch.username: "elastic"
elasticsearch.password: "P7cIq6DAmLi439kiclh"

# Kibana can also authenticate to Elasticsearch via "service account tokens".
# Service account tokens are Bearer style tokens that replace the traditional username/password based configuration.
# Use this token instead of a username/password.
# elasticsearch.serviceAccountToken: "my_token"

# Time in milliseconds to wait for Elasticsearch to respond to pings. Defaults to the value of
# the elasticsearch.requestTimeout setting.
#elasticsearch.pingTimeout: 1500

# Time in milliseconds to wait for responses from the back end or Elasticsearch. This value
# must be a positive integer.
#elasticsearch.requestTimeout: 30000

^G Help          ^O Write Out      ^F Where Is      ^K Cut           ^T Execute        ^C Location       M-U Undo        M-A Set Mark     M-] To Bracket
^X Exit         ^R Read File      ^\ Replace       ^U Paste         ^J Justify        ^/ Go To Line    M-E Redo        M-6 Copy        ^B Where Was
```

Or use `Elasticsearch.ServiceAccountToken` instead of username and password.

Generate `Elasticsearch.ServiceAccountToken` using this command

```
sudo /usr/share/elasticsearch/bin/elasticsearch-service-tokens create elastic/kibana kibana-token
```

```

GNU nano 8.6                               /etc/kibana/kibana.yml *

# The maximum payload size in bytes for incoming server requests.
#server.maxPayload: 1048576

# The Kibana server's name. This is used for display purposes.
# server.name: "kibana"

# ===== System: Kibana Server (Optional) =====
# Enables SSL and paths to the PEM-format SSL certificate and SSL key files, respectively.
# These settings enable SSL for outgoing requests from the Kibana server to the browser.
#server.ssl.enabled: false
#server.ssl.certificate: /path/to/your/server.crt
#server.ssl.key: /path/to/your/server.key

# ===== System: Elasticsearch =====
# The URLs of the Elasticsearch instances to use for all your queries.
elasticsearch.hosts: ["https://localhost:9200"]

# If your Elasticsearch is protected with basic authentication, these settings provide
# the username and password that the Kibana server uses to perform maintenance on the Kibana
# index at startup. Your Kibana users still need to authenticate with Elasticsearch, which
# is proxied through the Kibana server.
#elasticsearch.username: "elastic"
#elasticsearch.password: "P7cIqC6DAmLi439kiclh"

# Kibana can also authenticate to Elasticsearch via "service account tokens".
# Service account tokens are Bearer style tokens that replace the traditional username/password based configuration.
# Use this token instead of a username/password.
elasticsearch.serviceAccountToken: "AAEAAWVsYXN0aWMva2liYW5hL2tpYmFuYS10b2tlbjpSZmi1LVh6ZVNiyUhcRExnV2dHLTFn"

# Time in milliseconds to wait for Elasticsearch to respond to pings. Defaults to the value of
# the elasticsearch.requestTimeout setting.
#elasticsearch.pingTimeout: 1500

# Time in milliseconds to wait for responses from the back end or Elasticsearch. This value
Save modified buffer? [Y/N]
Y Yes          ^C Cancel

```

Add Kibana's encryption keys

Run this command to generate encryption keys

```
sudo /usr/share/kibana/bin/kibana-encryption-keys generate
```

```

└─$ sudo /usr/share/kibana/bin/kibana-encryption-keys generate
## Kibana Encryption Key Generation Utility

The 'generate' command guides you through the process of setting encryption keys for:
xpack.encryptedSavedObjects.encryptionKey
    Used to encrypt stored objects such as dashboards and visualizations
    https://www.elastic.co/guide/en/kibana/current/xpack-security-secure-saved-objects.html#xpack-security-secure-saved-objects

xpack.reporting.encryptionKey
    Used to encrypt saved reports
    https://www.elastic.co/guide/en/kibana/current/reporting-settings-kb.html#general-reporting-settings

xpack.security.encryptionKey
    Used to encrypt session information
    https://www.elastic.co/guide/en/kibana/current/security-settings-kb.html#security-session-and-cookie-settings

Already defined settings are ignored and can be regenerated using the --force flag. Check the documentation links for instructions on how to rotate encryption keys.
Definitions should be set in the kibana.yml used configure Kibana.

Settings:
xpack.encryptedSavedObjects.encryptionKey: 5de5fc29c0bc3fa876c9b0147fc0ea65
xpack.reporting.encryptionKey: 1301548fcffffe5e0e93520feb97911e
xpack.security.encryptionKey: 4d93cf0e7c255e2dbb88845c937e78b0

```

```

GNU nano 8.6                               /etc/kibana/kibana.yml *

# The path where Kibana stores persistent data not saved in Elasticsearch. Defaults to data
#path.data: data

# Specifies the path where Kibana creates the process ID file.
pid.file: /run/kibana/kibana.pid

# Set the interval in milliseconds to sample system and process performance
# metrics. Minimum is 100ms. Defaults to 5000ms.
#ops.interval: 5000

# Specifies locale to be used for all localizable strings, dates and number formats.
# Supported languages are the following: English (default) "en", Chinese "zh-CN", Japanese "ja-JP", French "fr-FR".
#i18n.locale: "en"
xpack.encryptedSavedObjects.encryptionKey: 5de5fc29c0bc3fa876c9b0147fc0ea65
xpack.reporting.encryptionKey: 1301548fcffffe5e0e93520feb97911e
xpack.security.encryptionKey: 4d93cf0e7c255e2dbb88845c937e78b0

# ===== Frequently used (Optional) =====
# ===== Saved Objects: Migrations =====
# Saved object migrations run at startup. If you run into migration-related issues, you might need to adjust these settings.
# The number of documents to migrate at once.
#ops.migration.batchSize: 1000

```

The elasticsearch uses **HTTPS**, So add **Elasticsearch's CA certificate** in kibana.yml

Go to **/etc/elasticsearch/certs** , From this folder copy the **http_ca.crt** file to kibana's folder or just specify it's path in kibana.yml. Here I'm copied and renamed it into **ca.crt**.

```
Session Actions Edit View Help
└─(root㉿kali)-[~/etc/elasticsearch/certs]
  └─# ls
    http_ca.crt  http.p12  transport.p12
  └─# cp http_ca.crt /etc/kibana/ca.crt
root@kali:~/etc/elasticsearch/certs
```

Add CA certificate to kibana.yml

elasticsearch.ssl.certificateAuthorities: ["path/to/your/ca.crt"]

elasticsearch.ssl.verificationMode: full

```
GNU nano 8.6                               /etc/kibana/kibana.yml *
#elasticsearch.compression: false

# List of Kibana client-side headers to send to Elasticsearch. To send *no* client-side
# headers, set this value to [] (an empty list).
#elasticsearch.requestHeadersWhitelist: [ authorization ]

# Header names and values that are sent to Elasticsearch. Any custom headers cannot be overwritten
# by client-side headers, regardless of the elasticsearch.requestHeadersWhitelist configuration.
#elasticsearch.customHeaders: {}

# Time in milliseconds for Elasticsearch to wait for responses from shards. Set to 0 to disable.
#elasticsearch.shardTimeout: 30000

# ===== System: Elasticsearch (Optional) =====
# These files are used to verify the identity of Kibana to Elasticsearch and are required when
# xpack.security.http.ssl.client_authentication in Elasticsearch is set to required.
#elasticsearch.ssl.certificate: /path/to/your/client.crt
#elasticsearch.ssl.key: /path/to/your/client.key

# Enables you to specify a path to the PEM file for the certificate
# authority for your Elasticsearch instance.
elasticsearch.ssl.certificateAuthorities: [ "/etc/kibana/ca.crt" ]]

# To disregard the validity of SSL certificates, change this setting's value to 'none'.
elasticsearch.ssl.verificationMode: full

# ===== System: Logging =====
# Set the value of this setting to off to suppress all logging output, or to debug to log everything. Defaults to 'info'
#logging.root.level: debug

# Enables you to specify a file where Kibana stores log output.
logging:
  appenders:
    file:
      type: file

^KG Help          ^O Write Out     ^F Where Is      ^K Cut           ^T Execute       ^C Location      M-U Undo        M-A Set Mark   M-] To Bracket
^X Exit          ^R Read File     ^N Replace      ^U Paste         ^J Justify       ^V Go To Line    M-B Redo        M-C Copy      M-B Where Was
```

Save and exit.

Why using CA certificate?

When Kibana connects to Elasticsearch over HTTPS, it must verify that the Elasticsearch server is legitimate, not an impostor or attacker pretending to be Elasticsearch.

Now check the status:

```

$ sudo systemctl status kibana
● kibana.service - Kibana
   Loaded: loaded -- since Sun 2025-10-12 15:38:15 UTC 15 days ago
     Active: active (running) since Sun 2025-10-12 15:38:15 UTC; 15 days ago
       Docs: https://www.elastic.co
 Main PID: 3497 (node)
 Tasks: 11 (limit: 18515)
 Memory: 1G (peak: 1.3G)
 CPU: 19.813s
 CGroup: /system.slice/kibana.service
           └─3497 /usr/share/kibana/bin/../node/glibc-217/bin/node /usr/share/kibana/bin/../src/cli/dist

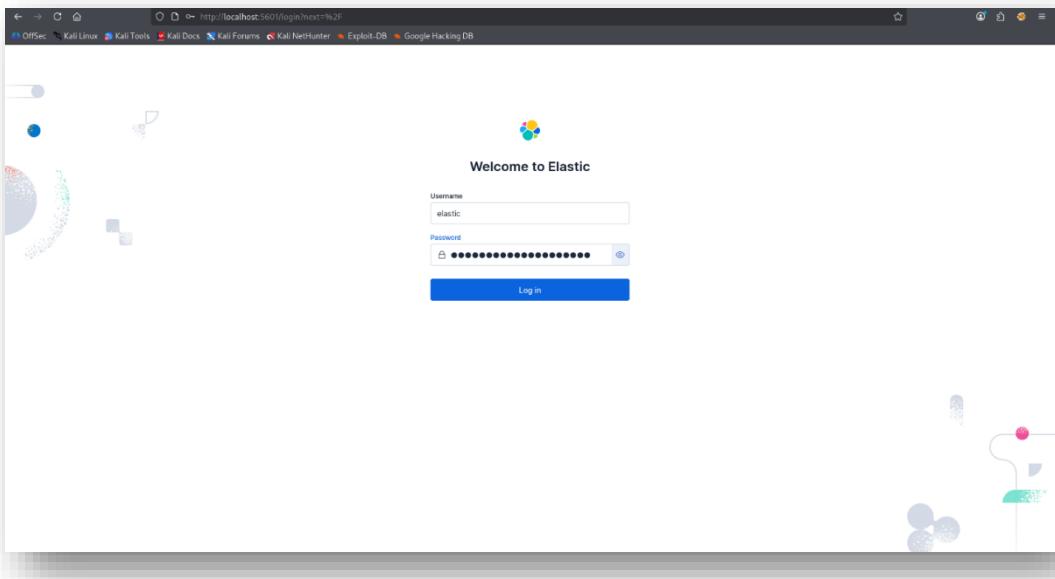
Oct 12 15:38:14 RA-7 systemd[1]: Started kibana.service - Kibana.
Oct 12 15:38:14 RA-7 kibana[3497]: {"log.level": "info", "@timestamp": "2025-10-12T10:08:15.632Z", "log.logger": "elastic-apm-node", "ecs.version": "8.10.0", "agentVersion": "4.13.0", "env": {"pid": 3497}, "type": "log", "tags": ["kibana"], "source": "http://localhost:5601/_status/stack", "offset": 15, "offset_end": 15}
Oct 12 15:38:15 RA-7 kibana[3497]: Native global console methods have been overridden in production environment.
[lines 1-15/15 (END)]

```

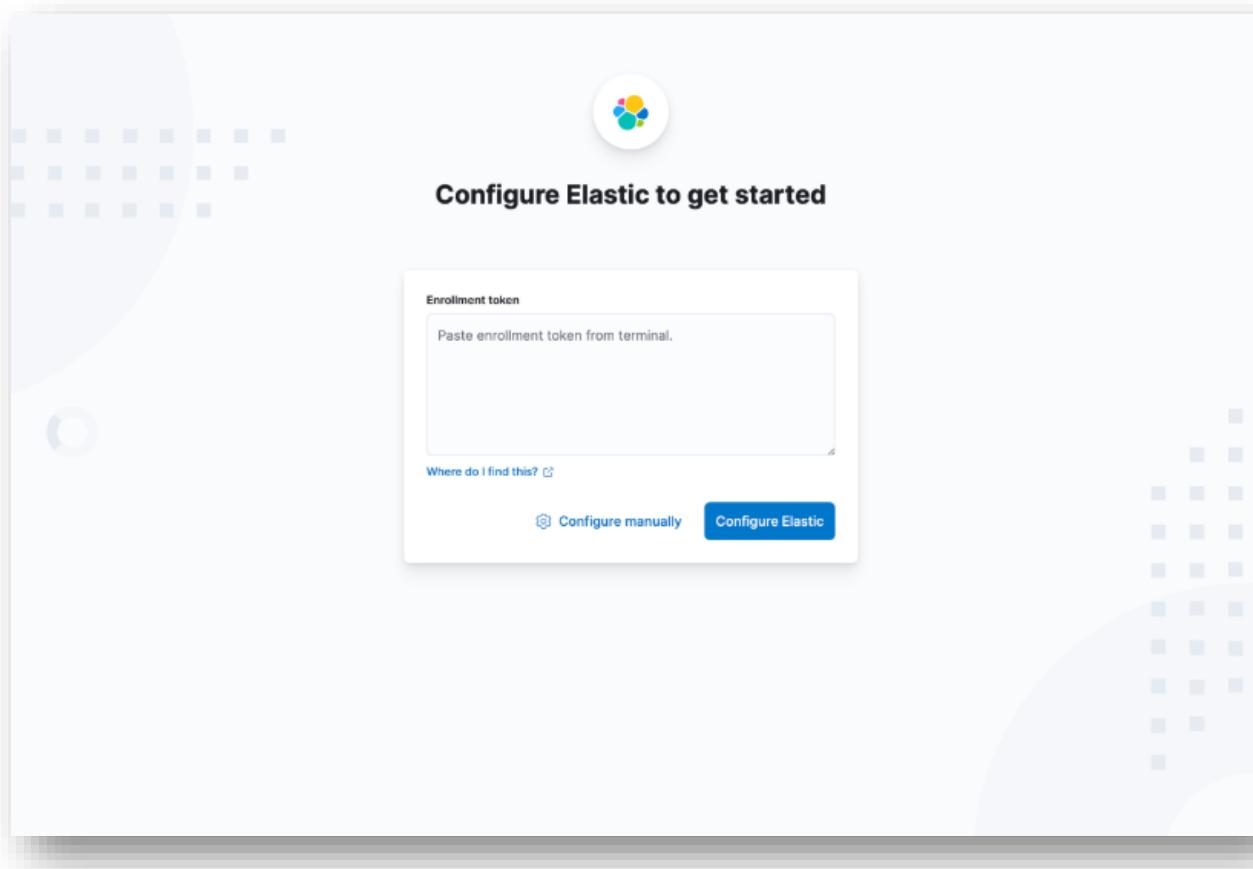
Then open your browser and go to:

<http://localhost:5061>

Login with your elasticsearch's username and password.



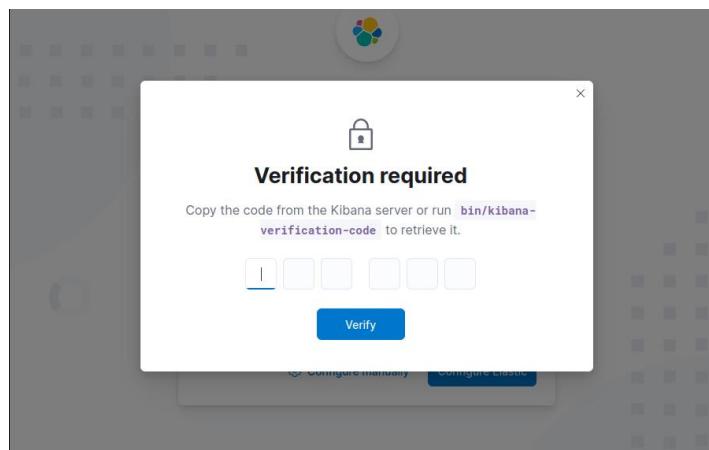
If Kibana asks for enrollment token:



Run this command to generate enrollment token:

```
sudo /usr/share/elasticsearch/bin/elasticsearch-create-enrollment-token -s kibana
```

after this kibana will ask for a verification code.



Get the verification code by this command:

```
sudo /usr/share/kibana/bin/kibana-verification-code
```

3) Configure Logstash

First make a pipeline configuration. This configuration file is the core instruction manual that tells Logstash which data to get (Input), how to clean it up (Filter), and where to send the final result (Output). Go to **/etc/logstash/conf.d** directory make one

```
sudo nano pipeline.conf
```

set the input and filter section like this

```
input {
  beats {
    port => 5044
  }
}

filter {
  # Suricata JSON processing
  if [message] and [message] =~ "^{.*}$" {
    json {
      source => "message"
      target => "suricata"
      skip_on_invalid_json => true
    }
    if [suricata][timestamp] {
      date {
        match => [ "[suricata][timestamp]", "ISO8601" ]
        target => "@timestamp"
      }
    }
  }

  if [suricata] {
    mutate {
      rename => { "[suricata][src_ip]" => "src_ip" }
      rename => { "[suricata][dest_ip]" => "dest_ip" }
      rename => { "[suricata][event_type]" => "event_type" }
      rename => { "[suricata][alert]" => "alert" }
    }
  }

  if [event_type] == "alert" or [suricata][event_type] == "alert" {
    mutate { add_tag => ["suricata-alert"] }
  }
}
```

Next add virustotal and alienvault integrations. Create account in this websites and get your api keys.

Virustotal: <https://www.virustotal.com/gui/user/codex07/apikey>

AlienVault OTX: <https://otx.alienvault.com/api>

Add this to filter section

```
if [suricata] {
  mutate {
    rename => { "[suricata][src_ip]" => "src_ip" }
    rename => { "[suricata][dest_ip]" => "dest_ip" }
    rename => { "[suricata][event_type]" => "event_type" }
    rename => { "[suricata][alert]" => "alert" }
  }
}

if [event_type] == "alert" or [suricata][event_type] == "alert" {
  mutate { add_tag => ["suricata-alert"] }
}

if [src_ip] {
  http {
    url => "https://otx.alienvault.com/api/v1/indicators/IPv4/%{src_ip}/general"
    verb => "GET"
    headers => {
      "X-OTX-API-KEY" => "2c5[REDACTED]cd3"
      "Accept" => "application/json"
    }
    target_body => "otx_response"
    connect_timeout => 5
    request_timeout => 20
    socket_timeout => 5
  }
}
```

^G Help ^O Write Out ^F Where Is ^K Cut ^T Execute ^C Location M-U Undo M-A Set Mark M-J To
^X Exit ^R Read File ^\ Replace ^U Paste ^J Justify ^/ Go To Line M-E Redo M-6 Copy ^B Whe

Virustotal api

```
if "_httprequestfailure" in [tags] or ![otx_response] {
    mutate { add_tag => ["otx-failed"] }
} else if [otx_response][pulse_info] and [otx_response][pulse_info][count] > 0 {
    mutate { add_tag => ["otx-malicious-ip"] }
}

if [dest_ip] {
    http {
        url => "https://www.virustotal.com/api/v3/ip_addresses/%{dest_ip}"
        verb => "GET"
        headers => {
            "x-apikey" => "1f2...773"
            "Accept" => "application/json"
        }
        target_body => "vt_response"
        connect_timeout => 5
        request_timeout => 20
        socket_timeout => 5
    }

    if "_httprequestfailure" in [tags] or ![vt_response] {
        mutate { add_tag => ["vt-failed"] }
    } else if [vt_response][data] and [vt_response][data][attributes][last_analysis_stats][malicious] > 0 {
        mutate { add_tag => ["vt-malicious-ip"] }
    }
}

# Auth.log parsing & filtering "Failed password"
if ![suricata] {
    if "Failed password" in [message] {
```

Prasing auth.logto filter Failed password. You can modify this by adding your custom filters.

```
}
```

```
# Auth.log parsing & filtering "Failed password"
if ![suricata] {
    if "Failed password" in [message] {
        grok {
            match => {
                "message" => "%{TIMESTAMP_ISO8601:iso_ts} %{HOSTNAME:host} sshd(%{POSINT:pid}\%): Failed password for (%:invalid user) ?%{USERNAME:user} from %{IP:src_ip} port %{INT:src_port} on %{IP:dst_ip} port %{INT:dst_port} at %{TIMESTAMP_ISO8601:timestamp} (%{WORD:auth_type})"
                tag_on_failure => ["_grokparsefailure_auth_fail"]
            }
        mutate {
            add_tag => ["ssh-auth-fail"]
        }
        if [iso_ts] {
            date {
                match => [ "iso_ts", "YYYY-MM-dd'T'HH:mm:ss.SSSSSSZ" ]
                target => "@timestamp"
            }
        } else {
            drop { }
        }
    }
}
```

Next set the output

Host => [“hostname/Ip_address:9200”]

Password => “your elasticsearch password”

ssl_certificateAuthorities => get this from /etc/elasticsearch/certs/http_ca.crt

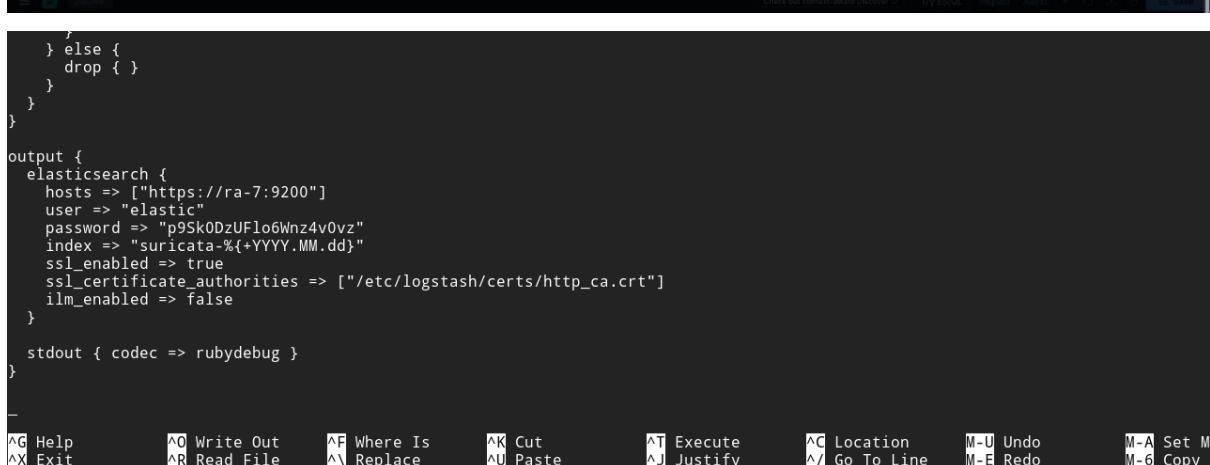


```

}
} else {
    drop { }
}
}

output {
    elasticsearch {
        hosts => ["https://ra-7:9200"]
        user => "elastic"
        password => "p9Sk0DzUFlo6Wnz4v0vz"
        index => "suricata-%{+YYYY.MM.dd}"
        ssl_enabled => true
        ssl_certificateAuthorities => ["/etc/logstash/certs/http_ca.crt"]
        ilm_enabled => false
    }
    stdout { codec => rubydebug }
}

-
```



Now edit logstash.yml

```
sudo nano /etc/logstash/logstash.yml
```

path.data: “/etc/lib/logstash”

```
GNU nano 8.6                                     logstash.yml *
```

```
# Settings File in YAML
#
# Settings can be specified either in hierarchical form, e.g.:
#   pipeline:
#     batch:
#       size: 125
#       delay: 5
#
# Or as flat keys:
#   pipeline.batch.size: 125
#   pipeline.batch.delay: 5
#
# _____ Node Identity _____
#
# use a descriptive name for the node:
# node.name: logstash
#
# If omitted the node name will default to the machine's host name
# _____ Data path _____
#
# Where directory should be used by logstash and its plugins
# for any persistent needs. Defaults to LOGSTASH_HOME/data
# path.data: "/var/lib/logstash"
# _____ Pipeline Settings _____
#
# The ID of the pipeline.
# pipeline.id: main
#
# Set the number of workers that will be parallelly execute the pipelines.
# pipeline.workers: 1
```

Set the pipeline.conf path

Path.config: “/etc/logstash/conf.d/filebeat-pipeline.conf”

```
GNU nano 8.6                                     logstash.yml *
```

```
# Disabled
# v0 (default)
# Pipelines defined before logstash 8 operated without ECS in mind. To ensure a
# single pipeline continues to operate as it did before your upgrade, opt-out
# of ECS for the individual pipeline in its "pipelines.yml" definition. Setting
# it here will set the default for _all_ pipelines, including new ones.
#
# pipeline.ecs_compatibility: v8
# _____ Pipeline Configuration Settings _____
#
# Where to fetch the pipeline configuration for the main pipeline
# path.config: "/etc/logstash/conf.d/filebeat-pipeline.conf"
# Pipeline configuration string for the main pipeline
# config.string:
# At startup, test if the configuration is valid and exit (dry run)
```

Edit the API Settings

api.enable: true

api.http.host: “0.0.0.0”

api.http.port: 9600-9700

```
GNU nano 8.6                                     logstash.yml *
```

```
# _____ API Settings _____
# Define settings related to the HTTP API here.
#
# The HTTP API is enabled by default. It can be disabled, but features that rely
# on it will not work as intended.
# api.enabled: true
#
# By default, the HTTP API is not secured and is therefore bound to only the
# host's loopback interface, ensuring that it is not accessible to the rest of
# the network. When secured with SSL and Basic Auth, the API is bound to _all_
# interfaces unless configured otherwise.
# api.http.host: "0.0.0.0"
#
# The HTTP API web server will listen on an available port from the given range.
# Values can be specified as a single port (e.g., 9600), or an inclusive range
# of ports (e.g., 9600-9700).
# api.http.port: 9600-9700
#
# The HTTP API includes a customizable "environment" value in its response,
# which can be configured here.
```

path.logs: “/var/log/logstash”

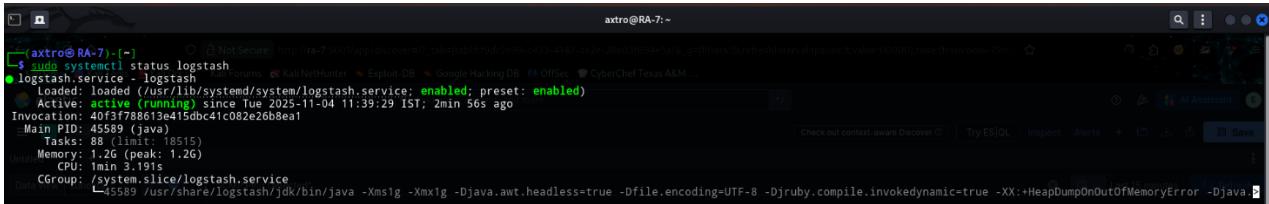
```
# path.dead_letter_queue:
# _____ Debugging Settings _____
#
# Options for log.level:
# * fatal
# * error
# * warn
# * info (default)
# * debug
# * trace
log.level: info
#
# Options for log.format:
# * plain (default)
# * json
#
# log.format: plain
# log.format.json.fix_duplicate_message_fields: true
#
path.logs: "/var/log/logstash"
# _____ Other Settings _____
#
# Allow or block running Logstash as superuser (default: false). Windows are excluded from the checking
# allow_superuser: false
```

Enable xpack.monitoring

```
# Defaults to heap, but can be switched to direct if you prefer using direct memory space instead.
# pipeline.buffer.type: heap
#
# ----- X-Pack Settings (not applicable for OSS build) -----
#
# X-Pack Monitoring
# https://www.elastic.co/guide/en/logstash/current/monitoring-logstash.html
# false allows the legacy internal monitoring (default: false)
#xpack.monitoring.allow.legacy.collection: false
#xpack.monitoring.enabled: true
#xpack.monitoring.elasticsearch.username: logstash_system
#xpack.monitoring.elasticsearch.password: ${password}
#xpack.monitoring.elasticsearch.proxy: ["https://proxy:port"]
#xpack.monitoring.elasticsearch.hosts: ["https://es1:9200", "https://es2:9200"]
# an alternative to hosts + username/password settings is to use cloud_id/cloud_auth
#xpack.monitoring.elasticsearch.cloud_id: logstash_system
#xpack.monitoring.elasticsearch.cloud_auth: logstash_system:password
# another authentication alternative is to use an Elasticsearch API key
#xpack.monitoring.elasticsearch.api_key: "id:api_key"
#xpack.monitoring.elasticsearch.ca_trusted_fingerprint: "xxxxxxxxxxxxxx"
#xpack.monitoring.elasticsearch.ssl.trusted_fingerprint: "xxxxxxxxxxxxxx"
#xpack.monitoring.elasticsearch.ssl.truststore.path: "/path/to/file"
#xpack.monitoring.elasticsearch.ssl.truststore.password: "password"
```

Save and exit

Check the status



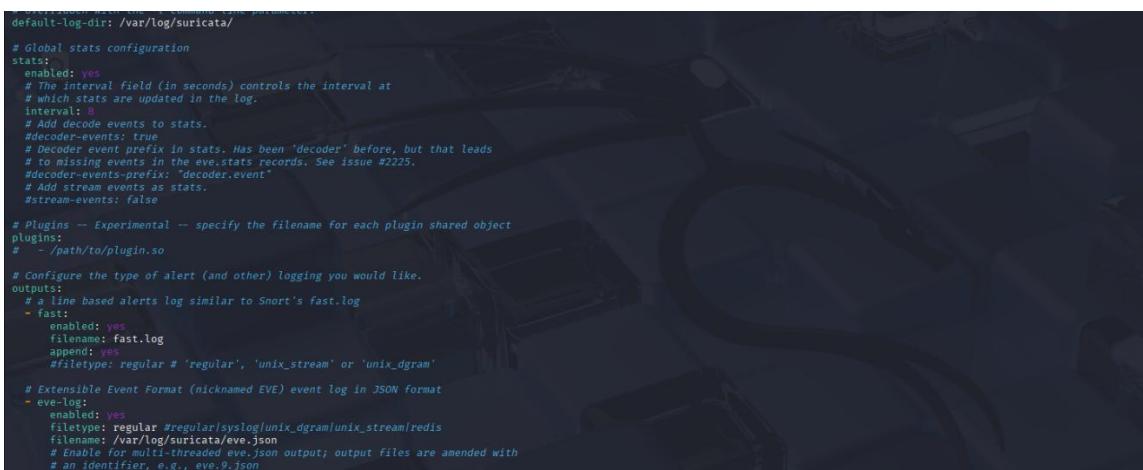
```
astro@RA-7:~$ sudo systemctl status logstash
● logstash.service - Logstash
   Loaded: loaded (/usr/lib/systemd/system/logstash.service; enabled; preset: enabled)
   Active: active (running) since Tue 2025-11-04 11:39:29 IST; 2min 56s ago
     Main PID: 45589 (java)
        Tasks: 88 (limit: 18515)
       Memory: 1.2G (peak: 1.2G)
          CPU: 1min 3.191s
         CGroup: /system.slice/logstash.service
           ↳ 45589 /usr/share/logstash/jdk/bin/java -Xms1g -Xmx1g -Djava.awt.headless=true -Dfile.encoding=UTF-8 -Druby.compile.invokedynamic=true -XX:+HeapDumpOnOutOfMemoryError -Djava
```

4) Configure Suricata

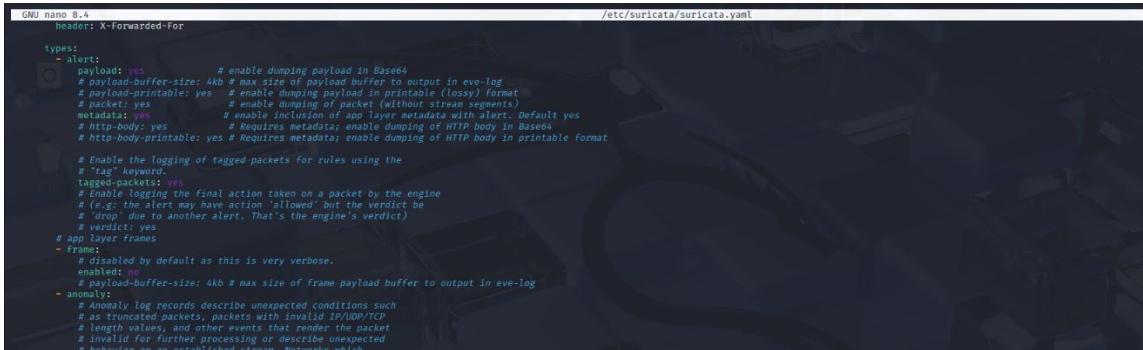
Go to `/etc/suricata` directory and edit `suricata.yaml` file

`sudo nano suricata.yaml`

set the output `eve.json`



```
default-log-dir: /var/log/suricata/
# Global stats configuration
stats:
  enabled: yes
# The interval field (in seconds) controls the interval at
# which stats are updated in the log.
interval: 8
# Add decode events to stats.
#decoder-events: true
# Decoder event prefix in stats. Has been 'decoder' before, but that leads
# to missing events in the eve.stats records. See issue #2225.
#decoder-events-prefix: decoder.event
# Add stream events as stats.
#stream-events: false
# Plugins -- Experimental -- specify the filename for each plugin shared object
plugins:
# - /path/to/plugin.so
# Configure the type of alert (and other) logging you would like.
outputs:
# a line based alerts log similar to Snort's fast.log
- fast:
    enabled: yes
    filename: fast.log
    append: yes
    #filetype: regular # 'regular', 'unix_stream' or 'unix_dgram'
# eve-log:
  enabled: yes
  #filetype: regular #regular|syslog|unix_dgram|unix_stream|redis
  filename: /var/log/suricata/eve.json
  # Enable for multi-threaded eve.json output; output files are amended with
  # an identifier, e.g., eve.9.json
```



```
GNU nano 8.4
header: X-Forwarded-For
types:
  types:
    alert:
      payload: yes      # enable dumping payload in Base64
      # payload-buffer-size: 4kb # max size of payload buffer to output in eve-log
      # payload-printable: yes  # enable dumping payload in printable (lossy) format
      # raw: yes             # enable dumping of packets without any processing
      metadata: yes        # enable including layer metadata with alert. Default yes
      # http-body: yes       # Requires metadata; enable dumping of HTTP body in Base64
      # http-body-printable: yes # Requires metadata; enable dumping of HTTP body in printable format
      # Enable the logging of tagged packets for rules using the
      # "tag" keyword
      tagged-packets: yes
      # Enable logging the final action taken on a packet by the engine
      # (e.g. the alert may have action 'allowed' but the verdict be
      # 'block'. This is another alert. That's the engine's verdict)
      # verdict: yes
      # app layer frames
      - frame:
          # disabled by default as this is very verbose.
          enabled: yes
          # payload-buffer-size: 4kb # max size of frame payload buffer to output in eve-log
      - anomaly:
          # Anomaly log records describe unexpected conditions such
          # as truncated packets, packets with invalid IP/UDP/TCP
          # headers, or other anomalies under the socket
          # threshold for further processing or describing
          # behavior on an established stream. Networks which
```

Now set your network interface to capture the traffic and log it in eve.json
here I'm using virtualbox, so mine is eth0

To know your interface:

ip link show OR ifconfig

then set it to yaml

```
# Linux high speed capture support
af-packet:
  - interface: eth0
  - interface: eth1
    # Number of receive threads. "auto" uses the number of cores
    threads: auto
    # Default clusterid. AF_PACKET will load balance packets based on flow.
    cluster-id: 99
    # Default AF_PACKET cluster type. AF_PACKET can load balance per flow or per hash.
    # This is only supported for Linux kernel > 3.1
    # possible value are:
    # * cluster_flow: all packets of a given flow are sent to the same socket
    # * cluster_cpu: all packets treated in kernel by a CPU are sent to the same socket
    # * cluster_qm: all packets linked by network card to a RSS queue are sent to the same
    #   socket. Required at least Linux 3.14.
    # * cluster_ebpf: eBPF file load balancing. See doc/userguide/capture-hardware/ebpf-xdp.rst for
    #   more info.
    # Recommended modes are cluster_flow on most boxes and cluster_cpu or cluster_qm on system
    # with capture card using RSS (requires cpu affinity tuning and system IRQ tuning)
    # cluster_rollback has been deprecated; if used, it'll be replaced with cluster_flow
    cluster-type: cluster_flow
    # In some fragmentation cases, the hash can not be computed. If "defrag" is set
    # to yes, the kernel will do the needed defragmentation before sending the packets.
    defrag: yes
    # To use the ring feature of AF_PACKET, set 'use-mmap' to yes
    #use-mmap: yes
```

Next set the rules for suricata

In the rules directory you can see many pre-defined rules but they are all commented and not defined in suricata.yaml . You can create custom rules for example local.rules. Here I made two rule files (nmap.rules and suricata.rules).

First set the rules file in suricata.yaml. Scroll down to the end

```
## Configure Suricata to load Suricata-Update managed rules.
##
default-rule-path: /etc/suricata/rules
rule-files:
  - suricata.rules
  - nmap.rules
##
## Auxiliary configuration files.
##
classification-file: /etc/suricata/classification.config
reference-config-file: /etc/suricata/reference.config
# threshold-file: /etc/suricata/threshold.config
##
## Include other configs
##
```

Save and exit, then check the status.

```
File Actions Edit View Help
(kali㉿kali)-[~/etc/network]
$ sudo systemctl enable suricata && sudo systemctl start suricata && sudo systemctl status suricata
Synchronizing state of suricata.service with SysV service script with /usr/lib/systemd/systemd-sysv-install.
Executing: /usr/lib/systemd/systemd-sysv-install enable suricata
● suricata.service - Suricata IDS/IDP daemon
  Loaded: loaded (/usr/lib/systemd/system/suricata.service; enabled; preset: disabled)
  Active: active (running) since Thu 2025-10-30 18:01:59 IST; 10min ago
  Invocation: 917b1ee6414a432ca1ad9a5da94d8382
    Docs: man:suricata(8)
           man:suricatasc(8)
           https://suricata.io/documentation/
  Main PID: 21508 (Suricata-Main)
    Tasks: 18 (limit: 4870)
   Memory: 88.2M (peak: 89.8M)
      CPU: 23.115s
     CGroup: /system.slice/suricata.service
             └─21508 /usr/bin/suricata -D --af-packet -c /etc/suricata/suricata.yaml --pidfile /run/suricata.pid

Oct 30 18:01:58 kali systemd[1]: Starting suricata.service - Suricata IDS/IDP daemon...
Oct 30 18:01:59 kali suricata[21508]: Info: conf-yaml-loader: Configuration node 'rule-files' redefined.
Oct 30 18:01:59 kali suricata[21508]: 1: suricata: This is Suricata version 7.0.11 RELEASE running in SYSTEM mode
Oct 30 18:01:59 kali systemd[1]: Started suricata.service - Suricata IDS/IDP daemon.

(kali㉿kali)-[~/etc/network]
```

After this make the rule files

Go to </etc/suricata/rules> directory.

`sudo nano suricata.rules`



```
GNU nano 8.4
suricata.rules *

# Alert ICMP Ping Detected
alert icmp any any -> any any (msg:"ICMP Ping Detected"; ip.proto:1; itype:8; sid:1000001; rev:1;)

# Alert SSH brute-force attempt
alert tcp any any -> any 22 (msg:"SSH brute-force attempt"; flow:to_server,established; content:"SSH"; depth:4; threshold:type threshold, track_by_src, count 5, seconds 60; classtype:attempted-admin; sid:1000002; rev:1;)
```

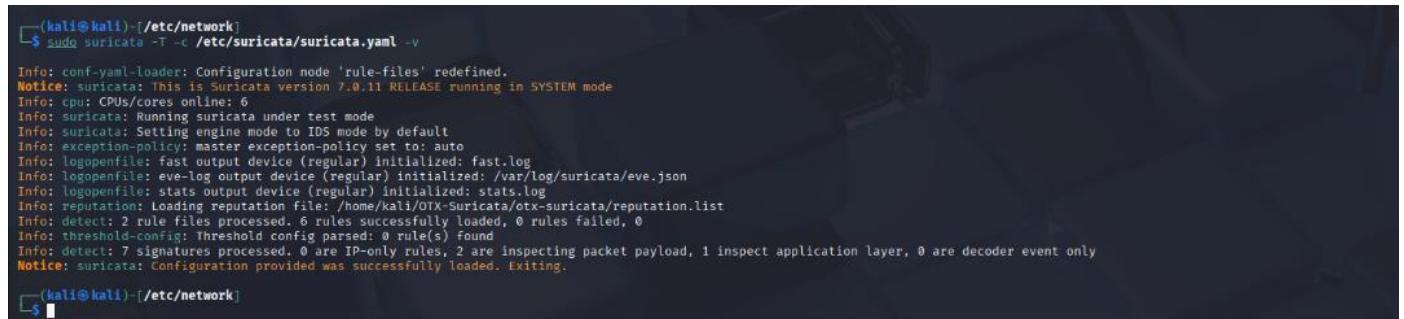
Here you can see two rules, The first one is to detect Ping attemps and the second one is to detect SSH brute-force attempt. Actually both this rules are for NIDS, that's why there is **any any → any any** which means from any source ip and source port to any destination ip and destination port. The second rule is to detect the login attempts to port 22 which is SSH service. Proper way to confirm this is to check the logs for “**Failed password**” from **auth.log**

If there are many logs with Failed password within a seconds, Then it's a Brute-Force attempt.

Like this create the nmap rules as well. Note that every rules sid must be unique, otherwise it won't work.

Run this command to check the rules are working properly or not:

`sudo suricata -T -c /etc/suricata/suricata.yaml -v`



```
(kali㉿kali)-[~/etc/network]
$ sudo suricata -T -c /etc/suricata/suricata.yaml -v
Info: conf-yaml-loader: Configuration node 'rule-files' redefined.
Notice: suricata: This is Suricata version 7.0.11 RELEASE running in SYSTEM mode
Info: cpu: CPUs/cores online: 6
Info: suricata: Running suricata under test mode
Info: suricata: Setting engine mode to IDS mode by default
Info: exception-policy: master exception-policy set to: auto
Info: logopenfile: fast output device (regular) initialized: fast.log
Info: logopenfile: eve-log output device (regular) initialized: /var/log/suricata/eve.json
Info: logopenfile: stats output device (regular) initialized: stats.log
Info: reputation: Loading reputation file: /home/kali/OTX-Suricata/otx-suricata/reputation.list
Info: detect: 2 rule files processed. 6 rules successfully loaded, 0 rules failed, 0
Info: threshold-config: Threshold config parsed: 0 rule(s) found
Info: detect: 7 signatures processed. 0 are IP-only rules, 2 are inspecting packet payload, 1 inspect application layer, 0 are decoder event only
Notice: suricata: Configuration provided was successfully loaded. Exiting.

(kali㉿kali)-[~/etc/network]
```

Check the eve.json is generated or not in the folder </var/log/suricata>

5) Configure Filebeat

Go to </etc/filebeat> directory and edit the filebeat.yml file

`sudo nano filebeat.yml`

go to the filebeat.input section and set the path to the logs which we want to analyse.

`/var/log/auth.log`

`/var/log/suricata/eve.json`

```

# Configuration file.

# ===== Filebeat inputs =====

filebeat.inputs:

# Each - is an input. Most options can be set at the input level, so
# you can use different inputs for various configurations.
# Below are the input-specific configurations.

# filestream is an input for collecting log messages from files.
- type: filestream

# Unique ID among all inputs, an ID is required.
id: my-filestream-id

# Change to true to enable this input configuration.
enabled: true

# Paths that should be crawled and fetched. Glob based paths.
paths:
- /var/log/auth.log
- /var/log/suricata/eve.log
#- /var/log/syslog
#- /var/log/boot.log
json.keys_under_root: true
json.add_error_key: true
#- c:\programdata\elasticsearch\logs\*

# Exclude lines. A list of regular expressions to match. It drops the lines that are

```

Next comment the output elasticsearch section and uncomment the output logstash section

hosts: [“your_IP/Host name:5044”] (Hostname of the system running Logstash)

It's better to give the hostname instead to IP. IP address will change if you connect to different network.

```

#ssl.certificateAuthorities: "/etc/filebeat/certs/http_ca.crt"  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
# ===== Logstash Output =====  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
output.logstash:
# The Logstash hosts  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
hosts: ["kali:5044"]  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
# Optional SSL. By default is off.  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
# List of root certificates for HTTPS server verifications  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
#ssl.certificateAuthorities: ['/etc/pki/root/ca.pem']  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
# Certificate for SSL client authentication  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
#ssl.certificate: "/etc/pki/client/cert.pem"  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
# Client Certificate Key  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
#ssl.key: "/etc/pki/client/cert.key"  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
# ===== Processors =====  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
processors:
- add_host_metadata:  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
  when.not.contains.tags: forwarded  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
- add_cloud_metadata: ~  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
- add_docker_metadata: ~  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
- add_kubernetes_metadata: ~  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
- add_fields:
  target: ''  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
  fields:  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This
  event_type: 'alert'  The decimal volume range for element 'range' (-14000 dB ~ -1000 dB) has negative maximum. This

```

Check the connection between filebeat and logstash. Note your logstash must be running otherwise it won't connect

sudo filebeat test output

```

(kali㉿kali)-[~/Pictures]
$ sudo filebeat test output
logstash: ra-7:5044 ...
connection...
parse host ... OK
dns lookup ... OK
addresses: 192.168.31.203
dial up ... OK
TLS... WARN secure connection disabled
talk to server ... OK

(kali㉿kali)-[~/Pictures]
$ 

```

Check for logs in kibana

Go to Kibana => Discover, change the index to suricata-* (If not click create one) OR you can simply use filebeat-* by changing the intex in pipeline's output setcion

Index => "filebeat-%{+YYYY.MM.DD}"

The screenshot shows the elastic App Discover interface. On the left, there's a sidebar with a list of fields: Data views, Default security data view, filebeat-* (with a dropdown menu), logs-* (selected and marked as Managed), metrics-* (also managed), Security alert data view, and suricata-* (also managed). A red arrow points from the text "To create" to the "Create a data view" button in the top right of the sidebar. Another red arrow points from the text "No results match your search criteria" to the "suricata-*" entry in the list.

To create

No results match your search criteria

Then you can see the logs

The screenshot shows the Elastic Stack interface with the following details:

- Header:** elastic, Find apps, content, and more, Try ES|QL, Inspect, Alerts, Save.
- Left sidebar:** Discover, Untitled, Data view, suricata-*, Search field names (with a count of 599).
- Main area:** A histogram showing event counts over time from November 17 to November 30, 2025. The x-axis is labeled "Auto interval" and "No breakdown".
- Table:** Documents [226] Field statistics, listing various fields and their values.
- Bottom navigation:** Rows per page: 100, Add a field, Last 15 minutes, Refresh, Date quick select.

Now check the suricata rules are triggered

Let's ping to the host system and search for the message that we set in the rule to detect ICMP ping.

```
[-] astro@RA-7:~$ ping 192.168.31.66
PING 192.168.31.66 (192.168.1.66) 56(84) bytes of data.
64 bytes from 192.168.31.66: icmp_seq=1 ttl=64 time=1.25 ms
64 bytes from 192.168.31.66: icmp_seq=2 ttl=64 time=1.12 ms
64 bytes from 192.168.31.66: icmp_seq=3 ttl=64 time=0.929 ms
64 bytes from 192.168.31.66: icmp_seq=4 ttl=64 time=1.35 ms
...
--- 192.168.31.66 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3005ms
rtt min/avg/max/mdev = 0.929/1.162/1.354/0.158 ms
[-] astro@RA-7:~$
```

The screenshot shows the Elastic Kibana Discover interface. The search bar at the top contains the query "ICMP Ping Detected". Below the search bar, the results are displayed under the heading "Documents (4)". The first document listed is a log entry from November 4, 2025, at 11:32:08.301. It contains several fields, including "event.original" which includes the alert message: "src_ip": "192.168.31.203", "dest_ip": "192.168.31.66", "proto": "ICMP", "type": "ping", "alert": "allowed", "signature_id": "000000", "rev": 1, "signature": "ICMP Ping Detected", "category": "", "severity": 3. A red box highlights this specific alert message.

Expand it

The screenshot shows the Kibana Log view for the expanded log entry. The log entry is displayed in JSON format. A red box highlights the "source_address" and "source_ip" fields, both of which are "192.168.31.203".

```

{
  "timestamp": "2025-11-04T11:32:08.301493+0530",
  "flow_id": "1549441992625333",
  "src_ip": "192.168.31.203",
  "dest_ip": "192.168.31.66",
  "proto": "ICMP",
  "type": "ping",
  "alert": "allowed",
  "signature_id": "000000",
  "rev": 1,
  "signature": "ICMP Ping Detected",
  "category": "",
  "severity": 3
}

```

The rule works properly and can see the alert message. Expand the log to get more details like source IP, source port, etc. You can set alerts in the free Kibana version, but your options are limited to the features available with the basic license. With the free version, you can create the rules themselves like index threshold or query rules, but you are restricted to using only a limited set of actions for notifications, primarily logging to an index. Sending notifications through more advanced channels like Slack, email, or webhooks requires a paid commercial license.

Dashboard Setup

1) Create dashboard

- Kibana → Analytics → Dashboards
- Click **Create dashboard**
- Click **Save**, then give a name for your dashboard and save it.

2) Create visualization

Lets create a visualization for ICMP Ping

- Analytics → visualize library → Create visualization → Lens.
- Choose your data view (filebeat-* or suricata-*)
- Set the visualization type to **Area**
- Add filter (+ button on the top left)

Field= suricata.eve.icmp_type

Operator= is

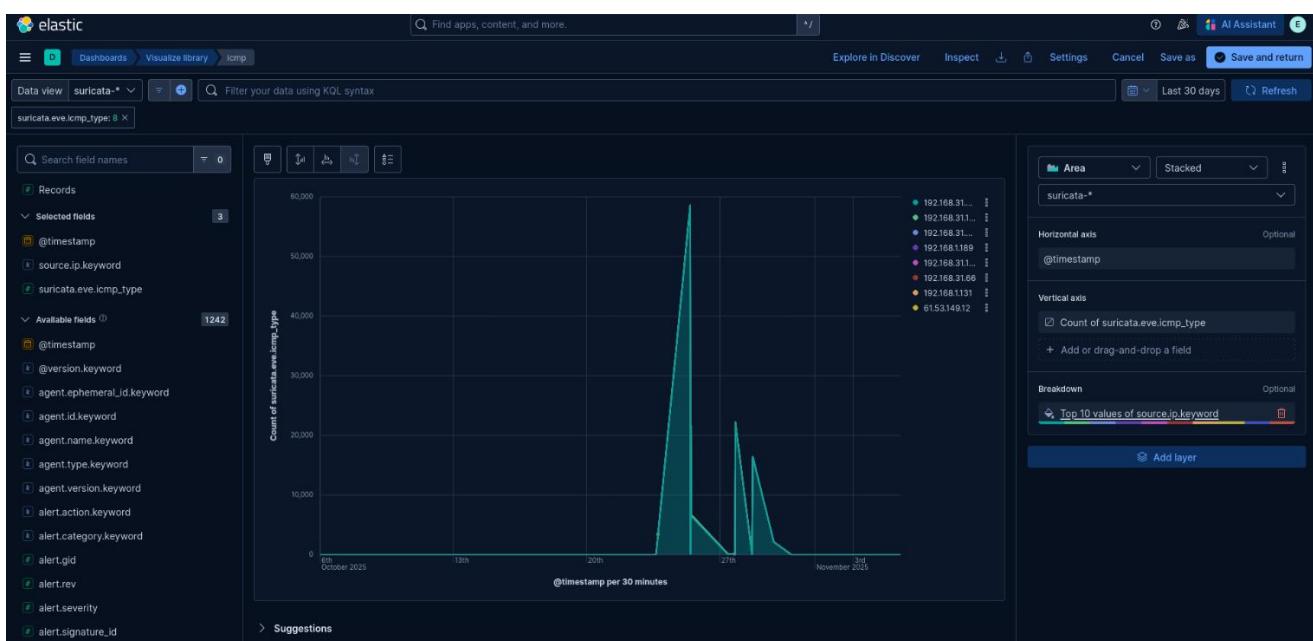
Field value= 8

- Drag **@timestamp** to the horizontal axis
- Click vertical axis and select **count** then set **suricata.eve.icmp_type** to the field.
- Breakdown

Field= source.ip.keyword (if this not working you can use any field related to source ip. Eg. src_ip)

No.of values= 10

- Click **Save** → Give name → Add to existing dashboard
- Search for your dashboard → Save → Go to dashboard



Visualization for Nmap scans

- Go to Lens
- Add filter : add filter based on the suricata rules we set earlier by their signature(msg).

Field= alert.signature (if this not working try **suricata.eve.alert.signature.keyword**).

Operator= is

Field value= msg from the suricata nmap rules (Eg. Nmap Xmas scan)

Click **OR** and do the same thing with other nmap signatures from the rules.

- Visualization type= **Pie**
- Slice by

Field= **suricata.eve.alert.signature.keyword** (Same as you select as the field for filter).

Add another field.

Field= **source.ip.keyword**

No.of volume= 10

Ranked by count of records

Visualization for SSH bruteforce attempt

First make sure your pipeline can parse logs from auth.log. If not, add a grok filter. This is something you should handle carefully. Add a grok filter for auth.log without affecting your eve.json filter by separate them properly in the pipeline.

Afther this

- Go to **Lens**
- Choose your data view
- In the top filter bar, enter the KQL query: **message : “Failed password”**
- Drag **@timestamp** to the horizontal axis
- Select **Count** in vertical axis
- In breakdown, set **src_ip** and set top 10
- Save it and add to your dashboard.

You can create many visualizations like this.



Conclusion

The development of a Host-Based Intrusion Detection System (HIDS) using ELK Stack and Suricata demonstrates the effectiveness of open-source tools in strengthening system security and monitoring host level threats. By integrating VirusTotal and AlienVault OTX, the system gains access to real time threat intelligence, enhancing its ability to detect and correlate malicious activities efficiently.

While this project focuses on a host level monitoring, Suricata is inherently designed as a Network Intrusion Detection and Prevention system (NIDS/NIPS). When deployed on a network gateway or connected to a SPAN (Switched Port Analyzer) or mirror port, Suricata can capture and analyse traffic across the network, providing comprehensive visibility into network-level attacks.

This project demonstrated how modern cybersecurity frameworks can integrate host-level log analysis and external threat intelligence to provide a comprehensive defence against evolving cyber threats. Its modular and scalable design also supports future enhancements, including machine learning-driven anomaly detection and automated incident response, establishing a strong foundation for advanced cybersecurity implementations.